

Metehan Yilmaz

metehan.yilmaz@email.com | +90 532 111 2233 | linkedin.com/in/metehanyilmaz | github.com/metehanyilmaz

PROFESSIONAL SUMMARY

Senior Full-Stack Software Engineer with 6 years of experience building scalable distributed systems and AI-integrated products. Passionate about clean architecture, LLM tooling, and open-source contributions.

EDUCATION

Middle East Technical University (METU) — B.Sc. Computer Engineering

2014 – 2018 | GPA: 3.71 / 4.00

LANGUAGES

- English: C2 — Proficient
- German: B2 — Upper Intermediate

WORK EXPERIENCE

Senior Software Engineer — Trendyol — Istanbul (2021 – Present)

Led the design of a real-time recommendation engine (Python, Kafka, Redis) serving 40M+ users. Integrated OpenAI API for product description enrichment. Mentored a team of 5 junior engineers.

Software Engineer — Getir — Istanbul (2018 – 2021)

Built courier assignment microservice in Go. Reduced delivery ETA estimation error by 18% using ML-based route prediction (scikit-learn, XGBoost).

PROJECTS

- LLM Document Analyzer:** FastAPI + LangChain + ChromaDB — RAG pipeline for legal document Q&A.; 200+ GitHub stars.
- OpenCV Face Attendance:** Real-time face recognition attendance system for university labs. PyTorch, MTCNN.
- Distributed Task Queue:** Custom Celery-like task queue in Python supporting Redis and RabbitMQ backends.

INTERNSHIPS

Software Development Intern — Aselsan — Ankara (Summer 2017)

Developed embedded C++ modules for radar signal processing unit.

Backend Intern — Arçelik R&D; (Summer 2016)

Built REST APIs for IoT device management dashboard using Django.

TECHNICAL SKILLS

Python, Go, FastAPI, Docker, Kubernetes, Kafka, PostgreSQL, Redis, PyTorch, TensorFlow, LangChain, OpenAI API, AWS

AI & MACHINE LEARNING

AI & Machine Learning: Built production RAG pipeline with LangChain and ChromaDB. Fine-tuned BERT for Turkish NLP tasks. Completed Stanford ML Specialization (Coursera). Contributor to Hugging Face Transformers library (3 merged PRs). Experience with OpenAI, Anthropic, and self-hosted Ollama deployments.